RASHMITHA J

Bengaluru, Karnataka, 560068

Phone: +919066747351 Email: rashmithajagadish1@gmail.com LinkedIn: linkedin.com/in/rashmitha-j-987765332

GitHub: github.com/rashmitha-j Portfolio: my-portfolio-site-ivory-alpha.vercel.app

EDUCATION

Don Bosco Institute Of Technology **Bengaluru, India**

Bachelor Of Engineering in Information Science and Engineering, CGPA 9.02 October 2024 - June 2028

SKILLS SUMMARY

- **Languages:** Python, C++, JavaScript, TypeScript, Java, HTML/CSS, SQL
- **Frameworks:** ReactJS, Next.js, Node.js, Express.js, FastAPI, Tailwind CSS
- **AI / LLM:** Google Gemini API, RAG, Embeddings, Vector Search, Prompt Engineering
- **Databases:** PostgreSQL (pgvector), MongoDB, MySQL
- **Others:** DSA, OOPS, DBMS, OS, CN
- **Tools & DevOps:** Git, GitHub, GitHub Actions (CI), pytest, Vercel, Render
- **Soft Skills:** Analytical Thinking, Problem Solving, Communication, Time Management

PROJECTS

Nexora AI – Document Q&A Chatbot with RAG | nexora-ai-jcdz.onrender.com October 2026

Tech Stack: Python, Chainlit, Google Gemini, PostgreSQL + pgvector (Neon), asyncpg, MetaGPT, FastAPI, pytest, GitHub Actions, Render

- Built a production-deployed **generative-AI assistant** that answers questions about uploaded PDF, TXT, MD, CSV and JSON files with **cited sources**, streaming answers token by token from Google Gemini.
- Implemented a **RAG pipeline**: text extraction, ~1,200-character chunking, **gemini-embedding-001** embeddings and cosine search in **pgvector** with an HNSW index, scoped so each user only searches their own documents.
- Kept the app running on free quotas with **automatic fallback across six Gemini models**, skipping exhausted or overloaded models and retrying only transient errors.
- Hardened the public endpoint with a **server-side model allow-list**, per-IP and site-wide rate limits, timeouts on every call and prompt-injection defences; added optional sign-in (PBKDF2) and saved chat history.
- Wrote **57 automated tests** with mocked LLM providers in GitHub Actions CI, and integrated a sandboxed **MetaGPT multi-agent** worker for a /build command.

DocuMind AI – Private PDF Question Answering | docu-mind-ai-delta.vercel.app October 2026

Tech Stack: Next.js, React, TypeScript, Tailwind CSS, Python, FastAPI, SQLAlchemy, PostgreSQL, LlamaIndex, Google Gemini, pytest

- Built a **full-stack RAG app** where users upload PDFs to a private library, select documents and get Gemini answers with **citations showing the document, page number and supporting passage**.
- Developed a **FastAPI** backend with Argon2id password hashing, hashed bearer-token sessions and per-user isolation on every route; the **Next.js** frontend is deployed on Vercel and the API on Render.
- Enforced per-user document and storage quotas inside a **database transaction with row locks**, so parallel uploads can't exceed limits; duplicate files reuse their existing index instead of re-embedding.
- Added **wall-clock deadlines** to every Gemini call with a fallback model, strict PDF validation, and rate limiting keyed on the real client IP behind Render's proxies.
- Wrote **80 offline pytest tests** with fake embeddings, run in GitHub Actions CI with frontend lint, type check and build.

CERTIFICATES

Complete Job-Ready Data Science Course – by Udemy

Learned data analysis, exploratory data analysis, machine learning fundamentals, and model building using Python.